

HotellIQ Data-Pipeline Documentation

Team Members:

- Yaksh Ajay Shah
 - Devarshi Anil Mahajan
 - Sarthak Vikas Sonawane
 - Rajkesh Prakash Shetty
 - Praveen Ramkumar
 - Rakshith Reddy Kokonda
-

Project Overview

The **HotellIQ-Data-Pipeline** is designed to extract, transform, and load hotel-specific data for an intelligent chatbot system. The pipeline processes hotel metadata, room information, amenities, policies, and reviews, transforming this data into a structured, query-optimized format in a PostgreSQL database.

Purpose: Extract hotel metadata and reviews, transform/validate, and load into a Postgres (Cloud SQL) database for the HotellIQ chatbot RAG system.

Stack:

- Apache Airflow (orchestration via docker-compose)
 - Python scripts (ETL logic)
 - PostgreSQL (Cloud SQL)
 - Google Cloud Storage (optional bucket utilities)
-

Contents of README Documentation

1. Prerequisites and Environment Setup
2. Steps to Run the Pipeline
3. Code Structure and Explanation
4. Data Flow and Architecture
5. Reproducibility and Data Versioning
6. Error Handling and Logging

1. Prerequisites and Environment Setup

Prerequisites

- **Operating System:** Linux machine (project tested on Linux)
- **Docker & Docker Compose:** For running Airflow locally
- **Python:** Version 3.9+ with pip
- **Database Access:** Credentials for Cloud SQL (PostgreSQL) or local Postgres
- **Google Cloud SDK (Optional):** For Cloud SQL Proxy

Installation Steps

Step 1: Clone the Repository

```
git clone https://github.com/Rakshith-Reddy-K/hotel-iq.git  
cd hotel-iq/data_pipeline
```

Explanation:

- Clones the HotelIQ repository from GitHub
- Navigates to the `data_pipeline` subdirectory where all pipeline code resides

Step 2: Install Python Dependencies

```
python -m pip install -r requirements.txt
```

Explanation:

- Installs required Python packages including:
 - `apache-airflow` - Workflow orchestration
 - `psycopg2-binary` - PostgreSQL database adapter
 - `pandas` - Data manipulation
 - `python-dotenv` - Environment variable management
 - Other dependencies for ETL operations

Step 3: Configure Environment Variables

Create a `.env` file in the `data_pipeline/` folder with the following variables:

```
# Database Configuration
DB_HOST=localhost      # Or Cloud SQL proxy host
DB_PORT=5432          # Or Cloud SQL proxy port
DB_NAME=hoteliq_db
DB_USER=your_db_username
DB_PASSWORD=your_db_password

# CSV Data Location
LOCAL_CSV_PATH=intermediate/csv

# (Optional) GCP Configuration
GCP_PROJECT_ID=your-project-id
GCS_BUCKET_NAME=hoteliq-data-bucket
# Gemini API Key
GEMINI_API_KEY = your_gemini_api_key
```

Explanation:

- **DB_variables***: Used by `db_pool.py` to establish database connections
- **LOCAL_CSV_PATH**: Directory where CSV files are located (defaults to `intermediate/csv`)
- **GCP variables**: For optional cloud storage integration via `bucket_util.py`

Step 4: (Optional) Start Cloud SQL Proxy

If using Google Cloud SQL, configure and start the Cloud SQL Proxy:

1. Edit `scripts/run_proxy.sh` with your:
 - Project ID
 - Region
 - Instance name
 - Service account credentials path
2. Run the proxy:
`bash scripts/run_proxy.sh`

Explanation:

- Cloud SQL Proxy creates a secure tunnel to your Cloud SQL instance
- The proxy runs on `localhost:5432` (or configured port)
- `db_pool.py` connects to the proxy, which forwards to Cloud SQL
- This enables secure access without whitelisting IPs

2. Steps to Run the Pipeline

Using Airflow (Recommended for Production)

Step 1: Start Docker Services

docker-compose up -d

Explanation:

- Starts Postgres, Redis, Airflow webserver, scheduler, and worker
- Runs in detached mode (`-d`) so services run in background
- Services defined in `docker-compose.yaml`

Step 2: Access Airflow Web UI

Navigate to: <http://localhost:8080>

Explanation:

- Airflow UI allows monitoring and triggering DAGs
- Default credentials typically: `airflow / airflow`
- DAG file: `dags/data_pipeline_airflow.py`

Step 3: Verify DAG Availability

Ensure the `data_pipeline` DAG appears in the Airflow UI:

- Check that `dags/` directory is properly mounted in `docker-compose.yaml`
- Refresh the DAG list if needed

Step 4: Prepare CSV Files

Ensure these CSV files exist in `intermediate/csv/`:

- `hotels.csv`
- `rooms.csv`
- `reviews.csv`
- `amenities.csv`
- `policies.csv`

Explanation:

- The loader expects these specific filenames
- Each CSV must have headers matching column definitions in `load_to_database.py`

- Empty cells should be empty strings (converted to NULL by loader)

Step 5: Trigger the DAG

In the Airflow UI:

1. Find the `data_pipeline` DAG
2. Toggle it to "ON"
3. Click "Trigger DAG" to start execution

Explanation:

- DAG orchestrates: Extract → Transform → Validate → Load
- Monitor task progress in the Airflow UI
- View logs for each task by clicking on task boxes

3. Code Structure and Explanation

Repository Structure

```

hotel-iq
├── README.md                      # Project documentation
└── data_pipeline
    ├── dags
    │   └── data_pipeline_airflow.py    # Airflow DAG definition
    ├── src
    │   ├── accumulated.py           # Utilities for accumulating or aggregating data
    │   ├── batch_selection.py        # Logic for selecting data batches for processing.
    │   ├── bucket_util.py           # GCS utilities (optional)
    │   ├── filtering.py             # Functions for filtering datasets
    │   ├── path.py                  # Path utilities for consistent file location handling.
    │   ├── prompt.py                # Prompt templates or helpers
    │   ├── state_management.py      # State tracking for batch or pipeline runs
    │   ├── transform.py             # Data transformation logic
    │   ├── utils.py                 # Util functions to extract hotel data
    │   └── utils_gemini.py          # Utilities for Gemini API or model integration.
    ├── docker-compose.yaml            # Airflow services configuration
    ├── requirements.txt
    └── scripts
        └── run_proxy.sh
    └── sql
        ├── db_pool.py               # Database connection pool
        ├── queries.py                # SQL DDL and query helpers
        └── test_connection.py        # Connection test script

```

3.1 Pipeline Flow

The data pipeline executes in 6 sequential stages with some parallel processing, orchestrated by the Airflow DAG in `data_pipeline_airflow.py`

Stage 1: Conditional Filtering

The pipeline starts by checking if city-filtered data already exists in GCS to avoid redundant processing

If filtering is needed:

1. Downloads raw hotels.txt from GCS and filters by city (e.g., "Boston")
2. Downloads raw reviews.txt and filters reviews for city hotels
3. Uploads filtered CSVs back to GCS

If filtering exists: Skips directly to batch selection

Both paths converge at a join point before proceeding.

Stage 2: Batch Selection

The pipeline selects the next batch of unprocessed hotels using state tracking:

Downloads state.json to identify already-processed hotel IDs

Filters out processed hotels and selects next 25 (configurable via BATCH_SIZE)

Saves batch_hotels.csv to GCS

Filters reviews for the current batch hotels and saves batch_reviews.csv

Stage 3: Parallel Enrichment

Two tasks run in parallel to enrich the batch data:

Path A - Ratings Computation:

1. Calculates aggregate ratings (mean overall, cleanliness, service, etc.) from batch reviews
2. Saves batch_ratings.csv

Path B - AI Enrichment:

1. Iterates through each hotel in the batch
2. Calls Gemini Live 2.5 Flash API to extract comprehensive hotel data (rooms, amenities, policies)
3. Includes 5-second delays between API calls for rate limiting
4. Saves enrichment data as JSONL

Stage 4: Data Transformation

After both enrichment paths complete, the pipeline merges all data:

1. Loads batch hotels, enrichment JSONL, and ratings CSV
2. Skips hotels with enrichment errors
3. Merges original hotel data with AI enrichment and ratings
4. Normalizes into 5 database-ready CSVs: batch_hotels.csv, batch_rooms.csv, batch_amenities.csv, batch_policies.csv, batch_reviews.csv

Stage 5: Database Load

The pipeline loads the normalized data into PostgreSQL:

1. Creates tables if they don't exist (hotels, rooms, reviews, amenities, policies)
2. Bulk inserts data using psycopg2.extras.execute_values()

Stage 6: State Management

Finally, the pipeline updates tracking state:

1. Appends batch results to accumulated city-wide datasets in GCS
2. Updates state.json with processed hotel IDs and batch metadata

3.2 Core Components Explanation

3.2.1. Data Processing

a. transform.py - Handles data transformation and aggregation

`compute_aggregate_ratings()` - Computes per-hotel mean ratings from reviews

`prepare_hotel_data_for_db()` - Merges batch hotels, enrichment data, and ratings into normalized CSVs for database loading

b. utils.py - Provides parsing, cleaning, and data manipulation utilities

`parse_raw_hotels()` - Parses JSONL hotel data into DataFrames

`parse_raw_reviews()` - Parses JSONL review data into DataFrames

`merge_hotel_data()` - Merges hotel metadata with Gemini enrichment data

Safe conversion functions (`safe_int()`, `safe_decimal()`, `safe_date()`, etc.) for type validation

3.2.2. Data Selection & Filtering

a. filtering.py - Filters raw data by city

`check_if_filtering_needed()` - Checks if filtered data already exists in GCS

`filter_all_city_hotels()` - Filters master hotels dataset for a specific city

`filter_all_city_reviews()` - Filters master reviews dataset for city hotels

b. batch_selection.py - Manages batch processing

`select_next_batch()` - Selects next batch of unprocessed hotels based on state tracking

`filter_reviews_for_batch()` - Filters reviews for hotels in current batch

3.2.3 AI Enrichment

a. prompt.py - Orchestrates hotel data enrichment using AI APIs

extract_hotel_data_gemini() - Extracts hotel data using Gemini API
enrich_hotels_gemini() - Enriches hotels in batch using Gemini

b. utils_gemini.py - Gemini API client implementation

get_hotel_data_async() - Async function to get hotel data from Gemini Live 2.5 Flash API
get_hotel_data() - Synchronous wrapper for async Gemini API calls

3.2.4 State & Storage Management

a. state_management.py - Tracks processing state

update_processing_state() - Updates state.json with processed hotel IDs and batch metadata to prevent duplicate processing

b. accumulated.py - Manages accumulated results

append_batch_to_accumulated() - Appends batch results to accumulated artifacts in GCS

c. bucket_util.py - GCS integration utilities

get_bucket() - Manages GCS client connection with singleton pattern
upload_file_to_gcs() - Uploads files to GCS with retry logic
download_file_from_gcs() - Downloads files from GCS

d. path.py - Centralized path management for local and GCS storage

Defines directory constants (RAW_DIR, FILTERED_DIR, PROCESSING_DIR, PROCESSED_DIR)
Provides path resolution functions for all data layers

3.3. SQL Core Files

a. db_pool.py - Connection Pool Management

This file implements a singleton connection pool using psycopg2.pool.ThreadedConnectionPool to efficiently manage database connections across concurrent Airflow tasks

Key functions:

initialize_pool() - Creates a connection pool with 1-20 connections using environment variables (CLOUD_DB_HOST, CLOUD_DB_PORT, etc.)

get_connection() - Context manager that safely acquires connections from the pool, handles commits/rollbacks, and returns connections automatically

close_all() - Closes all connections in the pool

The pool is initialized automatically on module import

b. queries.py - Database Schema and Bulk Loading

This file contains all SQL DDL statements and bulk insert logic.

Schema Creation Functions:

create_hotels_table() - Creates the main hotels table with 24 columns including ratings, location, and metadata

create_rooms_table() - Creates rooms table with foreign key to hotels and unique constraint on (hotel_id, room_type)

create_reviews_table() - Creates reviews table with unique constraint on (hotel_id, reviewer_name, review_date)

create_amenities_table() - Creates amenities table with unique constraint on (hotel_id, category)

create_policies_table() - Creates policies table with unique constraint on hotel_id

create_all_tables() - Orchestrates creation of all five tables in dependency order

Data Loading Functions:

bulk_insert_from_csvs(csv_dir) - Main entry point that reads 5 CSV files, groups data by hotel_id, and calls insert_one_hotel_complete() for each hotel

insert_one_hotel_complete() - Atomically inserts one hotel and all related records (rooms, amenities, policies, reviews) using psycopg2.extras.execute_values() for batch inserts

clean_dict_for_db() - Converts pandas NaN values to Python None for database compatibility

Utility Functions:

`list_tables()` - Queries information schema to list all tables in the public schema

All INSERT statements use ON CONFLICT DO NOTHING clauses for idempotency, allowing safe re-runs without duplicate key violations

c. test_connection.py - Connection Testing

Simple utility script to verify database connectivity and list tables

imports `list_tables()` from `queries.py` and can be run standalone to test the connection pool setup.


Usage in Pipeline

These SQL files are used by the Airflow DAG in two tasks:


1. `create_tables` task - Calls `create_all_tables()` to ensure schema exists
2. `insert_data_to_db` task - Calls `bulk_insert_from_csvs()` to load the 5 normalized CSV files produced by the transformation stage

4. Data Flow and Architecture

4.1 End-to-End Pipeline Flow



4.2 Airflow DAG Visualization



5. Explanation of Each Task

Task 1: check_if_filtering_needed

Type: BranchPythonOperator

Purpose: Determines whether city-based filtering of hotel and review data is required before proceeding with the pipeline.

Process:

- Evaluates whether filtered city data already exists in the intermediate storage
- Uses logic from `src.filtering.check_if_filtering_needed` module
- Makes a branching decision based on the presence/absence of filtered data
- Returns the task_id of the appropriate next task to execute

Output: Branches the DAG to either filtering tasks (`filter_all_city_hotels`) or skips directly to batch selection (`select_next_batch`), optimizing pipeline execution by avoiding redundant work.

Task 2: skip_filtering

Type: EmptyOperator

Purpose: Serves as a placeholder task when filtering operations are not required.

Process:

- Performs no operations
- Simply passes control to the next downstream task
- Maintains proper task dependency structure

Output: Continues DAG execution to the next logical step without any data processing.

Task 3: filter_all_city_hotels

Type: PythonOperator

Purpose: Filters the master hotel dataset to include only hotels from the specified target city.

Process:

- Reads the complete hotel dataset from source
- Applies city-based filtering criteria to identify relevant hotels
- Extracts only hotel records matching the target city
- Writes the filtered subset to an intermediate CSV file

Output: A filtered hotel data CSV containing only hotels from the specified city, reducing data volume for downstream processing.

Task 4: filter_all_city_reviews

Type: PythonOperator

Purpose: Filters the master reviews dataset to include only reviews associated with hotels in the specified city.

Process:

- Reads the complete reviews dataset from source
- Matches reviews to hotels from the filtered city hotel list
- Extracts only reviews for hotels in the target city
- Writes the filtered reviews to an intermediate CSV file

Output: A filtered reviews CSV containing only reviews for the city's hotels, ensuring relevance for downstream analytics.

Task 5: select_next_batch

Type: PythonOperator

Purpose: Selects the next batch of hotels for enrichment and processing, enabling incremental pipeline execution.

Process:

- Reads the filtered city hotel list
- Selects a configurable batch of hotels (e.g., 25 hotels) for processing
- Updates batch state tracking to maintain progress
- Identifies hotels not yet processed

Output: A batch hotel list file containing the subset of hotels to be enriched in this iteration, supporting scalable parallel processing.

Task 6: filter_reviews_for_batch

Type: PythonOperator

Purpose: Filters reviews to include only those corresponding to hotels in the current processing batch.

Process:

- Reads the city-level filtered reviews file
- Matches reviews to the hotels in the current batch
- Filters reviews based on batch hotel IDs

- Writes batch-specific reviews to an intermediate file

Output: A batch reviews CSV containing only reviews for the hotels being processed in the current batch, reducing data size for enrichment operations.

Task 7: `enrich_hotels_gemini`

Type: PythonOperator

Purpose: Enriches hotel data using the Gemini API or model to extract structured information such as amenities, policies, and other features.

Process:

- Reads batch hotel data and associated reviews
- Calls Gemini API/model with appropriate prompts
- Extracts structured information from unstructured text
- Parses and validates enrichment results

Output: Enriched hotel data file (JSONL or CSV format) containing extracted amenities, policies, and LLM-generated features for downstream use.

Task 8: `append_batch_to_accumulated`

Type: PythonOperator

Purpose: Appends the processed batch data to accumulated city-level datasets, building the complete dataset incrementally.

Process:

- Reads batch enrichment results
- Appends data to accumulated city-level files (hotels, amenities, policies)
- Performs deduplication to ensure data integrity
- Updates tracking metadata

Output: Updated accumulated files for the city, containing all processed batches to date, supporting recovery and incremental reruns.

Task 9: compute_aggregate_ratings

Type: PythonOperator

Purpose: Computes aggregate ratings and statistics for hotels based on filtered reviews.

Process:

- Reads city or batch reviews data
- Calculates overall ratings (average, median, etc.)
- Computes category-specific ratings (cleanliness, service, location, etc.)
- Aggregates sentiment and quality metrics
- Writes aggregate ratings to a CSV file

Output: Hotel ratings CSV containing summary statistics for analytics, reporting, and chatbot responses.

Task 10: prepare_hotel_data_for_db

Type: PythonOperator

Purpose: Prepares final normalized tables (hotels, rooms, amenities, policies) for efficient database loading.

Process:

- Merges hotel, enrichment, ratings, and review data
- Normalizes data into separate relational tables
- Formats data according to database schema requirements
- Writes final CSVs for each table

Output: Multiple CSV files (hotels.csv, rooms.csv, amenities.csv, policies.csv) ready for database import.

Task 11: create_all_tables

Type: PythonOperator

Purpose: Creates all required database tables if they do not already exist, ensuring schema readiness.

Process:

- Connects to the target database
- Executes DDL (Data Definition Language) statements
- Creates tables with appropriate schema, constraints, and indexes
- Handles idempotent creation (CREATE IF NOT EXISTS)

Output: Database tables ready for data loading, preventing load failures due to missing schema.

Task 12: bulk_insert_from_csvs

Type: PythonOperator

Purpose: Loads the prepared CSV files into the database using efficient bulk insert operations.

Process:

- Reads final prepared CSV files
- Establishes database connection
- Performs batch/bulk insert operations for each table
- Handles transaction management and error handling
- Logs insertion metrics (record counts, duration)

Output: Data successfully loaded into database tables, completing the ETL pipeline and making data available for analytics and chatbot queries.

6. Workflow Summary

Data Transformation Journey

Input: Raw hotel and review data files (CSV/Text)

Output: Normalized PostgreSQL database ready for analytics and chatbot queries


Transformations:

1. **Global → City-Specific:** All hotels/reviews → Target city data only
 - `filter_all_city_hotels`: Filter master hotel dataset by city
 - `filter_all_city_reviews`: Filter master reviews dataset by city
2. **Bulk → Batch:** City dataset → Manageable batches (25 hotels at a time)
 - `select_next_batch`: Select next hotel batch for processing
 - `filter_reviews_for_batch`: Extract reviews for current batch only
3. **Raw → Enriched:** Basic hotel data → AI-enhanced structured data
 - `enrich_hotels_gemini`: Extract amenities, policies, features via Gemini API
 - Transform unstructured text → Structured attributes
4. **Individual → Aggregated:** Hundreds of reviews → Hotel-level metrics
 - `compute_aggregate_ratings`: Calculate overall and category ratings
 - Generate summary statistics (average, sentiment scores)
5. **Scattered → Merged:** Separate enrichments → Unified hotel profiles
 - `prepare_hotel_data_for_db`: Merge ratings, enrichments, base data
 - Combine batch results into complete hotel records
6. **Denormalized → Normalized:** Flat files → Relational schema
 - Split into separate tables: hotels, rooms, amenities, policies, reviews
 - Establish proper foreign key relationships
7. **Incremental → Accumulated:** Batch results → City-wide dataset
 - `append_batch_to_accumulated`: Build up complete city data
 - Track processing state for recovery and monitoring
8. **Files → Database:** Prepared CSVs → PostgreSQL tables
 - `create_all_tables`: Establish database schema
 - `bulk_insert_from_csvs`: Load data efficiently into Cloud SQL

Pipeline Flow:

Raw Data → Filtered → Batched → Enriched → Aggregated → Normalized → Loaded → Accumulated

9. ER Diagram



9.1 HOTELS

Purpose: The main table storing core information about each hotel property.

Attributes:

- hotel_id (PK): Unique identifier for each hotel, used to link all other tables.
- official_name: The hotel's official name as displayed to guests.
- star_rating: Hotel classification from 1 to 5 stars, indicating quality level.
- address: Full street address of the hotel property.
- phone: Primary contact phone number for the hotel.
- email: Contact email address for guest inquiries.
- website: URL of the hotel's official website.
- overall_rating: Average rating calculated from all guest reviews (0-5 scale).
- total_reviews: Count of total reviews received by this hotel.
- cleanliness_rating: Average score specifically for room and property cleanliness.
- service_rating: Average score for staff service quality and responsiveness.
- location_rating: Average score for the hotel's location convenience.
- value_rating: Average score for price-to-quality ratio (value for money).
- year_opened: The year the hotel first opened for business.
- last_renovation: Date of the most recent property renovation or major update.
- total_rooms: Total number of guest rooms in the hotel.
- number_of_floors: Number of floors in the hotel building.
- additional_info: Extra details such as unique features, historical information, or enrichment data from external APIs.
- created_at: Timestamp when this hotel record was first created in the database.
- updated_at: Timestamp of the most recent update to this hotel record.

Relationships:

- Has (to ROOMS): One hotel can have multiple room types (one-to-many relationship).
- Has (to REVIEWS): One hotel can have multiple guest reviews (one-to-many relationship).
- Has (to AMENITIES): One hotel can have multiple amenities (one-to-many relationship).
- Has (to POLICIES): One hotel has one set of policies (one-to-one relationship, though stored as separate table).

9.2 ROOMS

Purpose: Stores information about different room types and configurations available at each hotel.

Attributes:

- **room_id** (PK): Unique identifier for each room type.
- **hotel_id** (FK): Links to the parent hotel in the HOTELS table, indicating which hotel this room belongs to.
- **room_type**: Category or classification of the room (e.g., "Standard Queen", "Deluxe King", "Executive Suite").
- **bed_configuration**: Description of bed arrangements in the room (e.g., "1 King Bed", "2 Queen Beds", "1 King + 1 Sofa Bed").
- **room_size_sqft**: Size of the room measured in square feet.
- **max_occupancy**: Maximum number of guests allowed in this room type.
- **price_range_min**: Minimum nightly rate for this room type (lowest price during off-season).
- **price_range_max**: Maximum nightly rate for this room type (highest price during peak season).
- **description**: Detailed description of room features, amenities, views, and other selling points.
- **created_at**: Timestamp when this room type was added to the database.

Relationships:

- **Belongs to (HOTELS)**: Each room type is associated with exactly one hotel via the **hotel_id** foreign key.

9.3 REVIEWS

Purpose: Stores individual guest reviews and ratings for hotels.

Attributes:

- **review_id** (PK): Unique identifier for each review.
- **hotel_id** (FK): Links to the hotel being reviewed in the HOTELS table.
- **overall_rating**: The reviewer's overall rating of the hotel (typically 1-5 scale).
- **review_text**: The full text of the guest's review, providing detailed feedback.
- **reviewer_name**: Name of the person who wrote the review (may be anonymized for privacy).
- **review_date**: Date when the review was posted or submitted.
- **source**: Where the review came from (e.g., "Google Reviews", "Booking.com", "TripAdvisor", "Direct").
- **created_at**: Timestamp when this review was added to the database.

Relationships:

- **Belongs to** (HOTELS): Each review is associated with exactly one hotel via the **hotel_id** foreign key.
- **Can Write**: Users can write multiple reviews for different hotels (implied relationship).

9.4 AMENITIES

Purpose: Stores information about facilities and services offered by each hotel.

Attributes:

- **amenity_id** (PK): Unique identifier for each amenity record.
- **hotel_id** (FK): Links to the hotel offering this amenity in the HOTELS table.
- **category**: Type or classification of the amenity (e.g., "WiFi", "Parking", "Dining", "Fitness", "Pool", "Business Center").
- **description**: Detailed description of the amenity, including specifics about what's offered.
- **details** (JSONB): Flexible structured data storing additional information such as:
 - Operating hours (e.g., {"hours": "6 AM - 10 PM"})
 - Pricing (e.g., {"cost": "Free"} or {"cost": "\$45/night"})
 - Specifications (e.g., {"speed": "100 Mbps"} for WiFi)
 - Location within hotel (e.g., {"location": "Rooftop, 12th floor"})

Relationships:

- **Belongs to** (HOTELS): Each amenity is associated with exactly one hotel via the **hotel_id** foreign key.

9.5 POLICIES

Purpose: Stores hotel policies, rules, and operational information that guests need to know.

Attributes:

- **policy_id** (PK): Unique identifier for each policy record.
- **hotel_id** (FK): Links to the hotel with these policies in the HOTELS table.
- **check_in_time**: Standard check-in time (e.g., "3:00 PM", "15:00").
- **check_out_time**: Standard check-out time (e.g., "11:00 AM", "11:00").
- **min_age_requirement**: Minimum age required to book a room (typically 18 or 21).
- **pet_policy**: Rules regarding pets, including fees, size restrictions, and allowed areas (e.g., "Pets allowed with \$75 fee. Maximum 2 pets, 50 lbs each").
- **smoking_policy**: Smoking rules and restrictions (e.g., "Non-smoking property. \$250 cleaning fee for violations").
- **children_policy**: Policies related to children, such as age limits for free stays (e.g., "Children under 12 stay free with existing bedding").
- **extra_person_policy**: Charges for additional guests beyond standard occupancy (e.g., "\$25 per person per night").
- **cancellation_policy**: Terms for canceling reservations, including deadlines and fees (e.g., "Free cancellation until 48 hours before check-in. Late cancellations charged one night").
- **updated_at**: Timestamp of the most recent update to these policies

Relationships:

- **Belongs to** (HOTELS): Each policy set is associated with exactly one hotel via the **hotel_id** foreign key.

10. Key Relationships Explained

Foreign Key Constraints

What is a Foreign Key? A foreign key is a column that creates a link between two tables. It ensures that data in one table matches data in another table.

In HotelIQ:

1. **ROOMS → HOTELS**
 - Each room's **hotel_id** must match a **hotel_id** that exists in the HOTELS table
 - This prevents creating rooms for non-existent hotels
 - If a hotel is deleted, all its rooms are automatically deleted (CASCADE DELETE)
2. **REVIEWS → HOTELS**
 - Each review's **hotel_id** must match a **hotel_id** in the HOTELS table
 - Ensures reviews are always connected to valid hotels
 - If a hotel is deleted, all its reviews are automatically deleted
3. **AMENITIES → HOTELS**
 - Each amenity's **hotel_id** must match a **hotel_id** in the HOTELS table
 - Prevents orphaned amenity records
 - If a hotel is deleted, all its amenities are automatically deleted
4. **POLICIES → HOTELS**
 - Each policy's **hotel_id** must match a **hotel_id** in the HOTELS table
 - Ensures policies are always tied to a specific hotel
 - If a hotel is deleted, its policies are automatically deleted

One-to-Many Relationships

HOTELS (1) —→ (Many) ROOMS

- **Meaning:** One hotel can have many different room types
- **Example:** Grand Boston Hotel has Standard Queen, Deluxe King, and Executive Suite rooms
- **Real-world:** A single hotel property offers multiple room options for guests

HOTELS (1) —→ (Many) REVIEWS

- **Meaning:** One hotel can receive many guest reviews
- **Example:** Grand Boston Hotel has 1,203 reviews from different guests
- **Real-world:** Multiple guests write reviews about their stay at the same hotel

HOTELS (1) —→ (Many) AMENITIES

- **Meaning:** One hotel can offer many different amenities
- **Example:** Grand Boston Hotel offers WiFi, Parking, Pool, Gym, Restaurant, etc.
- **Real-world:** Hotels provide various facilities and services to guests

HOTELS (1) —→ (One) POLICIES

- **Meaning:** One hotel has one set of policies (though stored as separate record)
- **Example:** Grand Boston Hotel has specific check-in/out times, pet rules, and cancellation terms
- **Real-world:** Each hotel has its own rules and operational guidelines

11. Summary

The HotelIQ Data Pipeline provides a robust, production-ready ETL system for hotel chatbot data. Key features include:

Technical Highlights

- **Batch UPSERT:** Efficient INSERT...ON CONFLICT for idempotent loads
- **Automatic Type Casting:** Converts CSV strings to proper database types
- **Connection Pooling:** Reuses database connections for performance
- **Foreign Key Safety:** Enforces load order to prevent constraint violations
- **Comprehensive Logging:** Tracks progress and errors at every step

Operational Benefits

- **Airflow Orchestration:** Automated scheduling and retry logic
- **Cloud SQL Integration:** Secure access via Cloud SQL Proxy
- **Reproducible:** Clear documentation and environment setup
- **Scalable:** Handles growing datasets with batch processing
- **Maintainable:** Modular code structure with single-responsibility functions

Data Handling

- **Transformation:** City filtering, aggregation, enrichment
- **Normalization:** Proper database schema with referential integrity
- **Error Handling:** Graceful failure recovery with detailed error messages